Instrument Transposition Literature Review – CSES Innovate (ORCA)

Ashish Bamba, Rahul Punji

1. Introduction

Instrument transposition (IT) – the process of adapting the musical material contents (i.e. MIDI) of one instrument to that of another – has long been a central element in music practice.

To the arranger who wants to swap a saxophone phrase onto an oboe, or reassign a piano right-hand figure to a string ensemble, it's no secret that having to manually modify pitch range, register, texture, and many more instrument features to ensure that a piece is playable on a target instrument is a tedious exercise.

Our purpose in this paper is to integrate enhanced practices to IT. Up until now, computational models have attempted to automate IT through limited actions, such as shifting pitches or swapping MIDI instrument labels [3][4]. In this literature, we explore the problem with enhanced methods that combine symbolic music processing and machine learning [3][4][6]. Beyond just moving notes, our motivation is to optimize the computational solution to IT by adapting the musical contents to the characteristics and constraints of the target instrument in real-time.

2. Instrument Transposition as Symbolic Transformation

2.1 Pitch-Centric Transposition

The pitch-centric transposition has been the standard approach to IT. This simple computational notion adds or subtracts a fixed number of semitone values to every note in its global pitch field. Without requiring any general knowledge of the target instrument, it can preserve rhythmic structure, harmonic intervals, and melodic contour.

This methodology may strike the ordinary audience as appealing on the surface because of its convenience.

However, this convenient solution to IT later proves to be suboptimal because it neglects core differences like comfortable tessitura, timbral response across registers, and physical constraints [2][5]. Oftentimes, the resulting music is just unplayable altogether at worst and musically awkward at best.

This limitation inspired us to explore more sophisticated techniques that treat pitch as only a one dimension feature out of countless more, instead of the sole determinant of musical material outputs.

2.2 MIDI-Level Instrument Remapping

An alternative yet suboptimal solution to IT is the MIDI-level Instrument Remapping approach. In this abstraction, we manage to maintain the exact same musical notes and instead switch the instrument patch (i.e. MIDI program number). Afterwards, we let the sequential stage (the audio synthesis process) update

the timbre contents, whether it be through SoundFonts, physical modeling, or neural audio synthesis [3][4].

The appeal is self-evident: It's oversimplified and convenient, arguably even more so than pitch-centric transposition. Unfortunately, there is one major downside: This solution requires instrument-agnostic musical contents. More so than not, the musical content inputs fail to meet this criteria. The solution can not withstand feasibility if one wanted to render polyphonic piano textures as monophonic wind instruments, or wide register passages reassigned to narrow range instruments [2][5].

Its shortcomings further highlight the importance of symbolic adaptation within the domain of IT itself and prior to synthesis.

3. Constraint-Aware Instrument Transposition

3.1 Range and Tessitura Constraints

A core concept that we will integrate in our IT literature is the explicit modeling of instrument range and tessitura.

A major conceptual advancement in the literature is the explicit modeling of instrument range and tessitura [2][5]. We will be careful to not hastily assume all pitches are equally viable, but instead we will reroute constrain aware systems to encode minimum, maximum, and preferred playable pitches for each instrument.

Here are concrete implementations of this aforementioned design philosophy:

1.) **Range aware transposition algorithms:** Instead of performing individual per-note tweaks, these find a key/octave that allow all notes to fit the target instrument's range [2].
2.) **Tessitura aware approaches** go further, preferring solutions that center musical material in comfortable registers rather than at extremes [5].

This design philosophy and its implementations have transformative implications. It longer approaches the IT solution through a series of local optimizations. Instead, we hunt for a global solution altogether by taking into consideration the entire distribution of pitches.

3.2 Playability as a Hard Constraint

A crucial component that is as equally important as the quality of the musical contents output is its playability. This has different implications in the context of different instrument types:

1.) For input contents where the instrument is monophonic, it is important that specific quality is kept intact in the output musical contents. This means ensuring each part only contains one note at a time, all of which are within the instrument's capabilities [2].

2.) For polyphonic contexts, it is important to take human constraints into consideration: hand span, fingering complexity, or breath constraints [1][5].

One of our motivations is to provide a strong guarantee that: if a solution is indeed found, then not only is it optimized for multi-dimensional features, but it is mechanically playable.

4. Instrument Transposition as Arrangement and Re Orchestration

4.1 Voice Assignment and Role Preservation

In this section, we will explore a special view of Instrument Transposition as being a case of Automatic Arrangement (AMA).

Through this framing, the system doesn't just tweak notes, but it also reasons about voices and parts (e.g. melody, harmony, or bass function) and can decide which target instruments are most suitable by determining which preserve most of the quality attributes [1][5][6]. Subsequently, voice aware systems may permute parts, apply octave displacements, or redistribute material across instruments to maintain structural clarity [2][5].

Ultimately, this piece acknowledges that musical function can not be treated as instrument-agnostic because our hypothesized model will produce better results for some target instruments over others.

4.2 Texture Preservation and Reduction

A recurring challenge in instrument transposition is texture mismatch – how many notes or lines can an instrument actually handle? An important discrepancy to note between polyphonic instruments and monophonic instruments is that the former can articulate dense chords while the latter can not. This precisely underscores why texture reduction is critical when transposing music from polyphonic to monophonic, just as is texture expansion when performing IT in the opposite direction [1][5].

Although it remains largely heuristic, the broader literature suggests that texture adaptation is inseparable if realism is to be an IT end-goal [6].

5. Learning-Based Approaches and Instrument Awareness

5.1 Neural Models for Instrumentation and Orchestration

Recent ML solutions to IT emphasize a standalone sequence to sequence learning model that takes in one instrument's symbolic parts and outputs the corresponding target instrument version [3][4][6].

While neural approaches excel at modeling subtle correlations that are difficult to encode as rules, purely data driven models may not suffice to guarantee playability or constraint satisfaction, producing suboptimal results that may be musically awkward [2][3].

This tension serves as one of our motivations to model a hybrid system that couples a rule-based constraint enforcement followed by a sequence-to-sequence (RNN) neural network architecture [3][4].

5.2 Instrument Identity as Structured Conditioning

Another crucial design principle we aim to consolidate is to decompose the instrumental inputs into structured attributes (i.e. range, tessitura, articulation set, and dynamic response) [6][7].

Conditioning models to operate on these attributes enables more flexible and generalizable instrument transposition, including adaptation to instruments not seen during training [6]. This perspective aligns closely with symbolic first pipelines, where instrument profiles can be explicitly encoded and used to guide both deterministic transformations and neural models.

6. Future Directions for Instrument Transposition in ORCA

The future of transposition in ORCA can accomplish a whole lot more than symbolic substitution and operate as a musical transformation process [3][4]. To achieve this, rather than thinking about how notes might need to be mechanically transposed from one instrument to another, ORCA should frame transposition as identifying a piece with a realization most analogous to a piece under the expressive capability of a certain musical instrument [5][6]. As far as human transposition service in musical instruments is concerned, this aligns with carrying out transposition in musical performance under a set of musical compromises [1]. In short, transposition can be an optimization problem to keep all musical elements intact, retain physical sound, allow for musical expression, and most importantly, make no distortions in melody, rhythm, and musical form at all [2][6]. In our pipeline, this is attempted in a step by step fashion with each step focusing on a different musical attribute and improving each time [6].

The most basic form of this setup will include an emphasis on using continual musical instrument descriptions [6][7]. Musical instrument descriptions will include other elements such as musical tessitura, possible articulation production, dynamic musical production capability, capability to produce a multiphonic sound, and possible limits in musical instrumentation [2][5]. Usefulness in musical instrumentation descriptions gives a chance to enforce a strict musical constraint and generalize this to unseen musical instruments [6][7]. Most importantly, these descriptions separate musical consideration from simple instrumentation in datasets and do not need instrumentation descriptions in musical processing [6]. Moreover, constraint enforcing preprocessing will be a necessity in musical processing before machine learning initiation since it will have to enforce basic musical instrumentation capability beforehand [2][5]. These musical normalizations ensure all musical elements can be played by a musical machine [2]. Furthermore, it will check all musical feasibility with a musical texture check to restrict

multiphonic musical instrumentation in case it is a monophonic musical machine and check musical instrumentation densities if it is a musical machine with multiphonic musical instrumentation capability [1][5]. At this stage of musical transposition, the output of this musical stage will feasible but not yet an idiomatic result [2][5].

Musical structure awareness and musical role will come into focus next. The primary goal of transposing an instrument is to maintain music roles by identifying them such as melody, accompanying roles, bass, and decorative lines, and musical structures such as phrases and cadences [1][5]. Through role-preserving transformations, it is possible to allow prioritization of melodic content, transform and simplify accompanying elements, and make decisions based on articulations and registers in a way that respects music rather than considering all occurrences equal [2][5][6]. With feasibility and roles determined, learned models will be able to take over to accomplish idiomatic detail [3][4][6]. Symbolic sequence-to-sequence models will be used to make decisions on register manipulation, articulations, rhythmic accents, and phrase settings all within the constrained solution space [3][4]. In addition to this, role information and register information will promote a better pathway for generalization without considering catastrophic failures [6].

Texture transpositioning will be considered a first-class problem rather than an after-effect by this system [1][5][6]. In regards to transposing polyphonic to monophonic or monophonic to polyphonic instruments, making a specific reduction in regards to chords, creating an arpeggiated strategy, or a time manipulation strategy based on redistribution of harmonic tones, or controlled expansion can be good strategies [1][5]. An explicit form of transposition allowing modeling, evaluation, and musical coherence in transposition will be important [6]. A strict boundary of separation between transposing an instrument symbolically and timbre generation will be the approach we take in this case [3][4]. Alternatively, decisions in music will have to be finalized at a symbolic level before synthesis takes place; this will not degrade integrity in regards to musical validity irrespective of technology used in synthesis [3][4].

Furthermore, in regards to evaluation and interaction, we will have to represent the challenge in a multicolinear way [6]. Quality in musical transpositioning will be measured and evaluated in regards to playability, musicality, validity in relationships with structures, and sound expressive validity; this can be accomplished through automatic methods and subjective evaluation on human judgment [1][5][6]. Human interaction methods such as dealing with registers in a locked manner, validity in texture reduction, or making decisions based on stylistic preferences will keep the system up to date with realities in creative work [5][6]. Through all these pathways in this constraint strategy, this system will move beyond just timbre substitution towards true instrument transposition [3][4][6].

7. SWOT Analysis

7.1 Strengths

1. Clear multi-stage design that mirrors how human arrangers think: constraints first, expression second [5].
2. Explicit instrument profiles give the system structure, interpretability, and strong generalization potential [6][7].

3. Separation of symbolic transposition from synthesis keeps the musical logic clean and reusable [3][4].
4. Role-aware and texture-aware processing boosts musical fidelity beyond basic pitch shifting [1][5].

7.2 Weaknesses

1. **Balancing constraints and expressivity**: Rule-based systems ensure playability but lack nuance; neural systems capture style but lack guarantees [3][4].
2. **Generalization across instruments**: Many models are tightly coupled to specific instrument sets or datasets [2][6].
3. **Texture adaptation**: Deciding when and how to reduce or expand texture remains largely heuristic [1][5][6].
4. **Evaluation**: There is no widely accepted metric for "good" instrument transposition, and human evaluation is costly [3][5].

7.3 Opportunities

1. Build a true multi-stage, musically aware transposition pipeline rather than simple pitch shifts or timbre swaps [3][4][6].
2. Leverage explicit instrument profiles to generalize to unseen instruments and support custom definitions [6][7].
3. Use constraint-first preprocessing + learned idiomatic refinement to deliver playability and musical authenticity [2][3][5].
4. Add human-in-the-loop controls for targeted adjustments, making the system usable in real arranging workflows.

7.4 Threats

1. Bad or incomplete instrument profiles can tank output quality across the entire pipeline [6][7].
2. Over-reliance on learned models risks unplayable or stylistically off outputs if constraints slip [2][3].
3. High complexity makes real-time or large-scale deployment harder compared to simple timbre substitution [3][4].

8. Conclusion

Instrument transposition occupies a space at the intersection of music theory, real-world constraints, and neural network architectures.

Early systems have attempted to reduce this problem to pitch shifting or timbre swapping to present a convenient solution, but these methods have proven to be suboptimal and too one-dimensional [2][3][4].

Our literature recognizes fundamental complexities that must be intact throughout our transposition pipeline but have previously been ignored in oversimplified approaches.

By building on top of pre-existing solutions, we hypothesize an exploration of a forward path that optimizes the solution to IT: A hybrid model containing a rule-based constraint filter followed by an RNN architecture, acting together as one unified framework [3][4][6].

9. References

[1] Chiu, Y.-H, M.-K., & Chen, Y.-C. 2008. *Automatic System for the Arrangement of Piano Reductions*. IEEE Transactions on Multimedia, 10(5), 843–856.

[2] McCloskey, Matthew, Gabrielle Curcio, Amulya Badineni, Kevin McGrath & Dimitris Papamichail. 2023.
Automated Arrangements of Multi-Part Music for Sets of Monophonic Instruments.
arXiv:2301.12084 [cs.SD].

[3] Brunner, G., N. C. (et al.). 2018.
MIDI-VAE: Modeling Dynamics and Instrumentation with Variational Autoencoders for Symbolic Music. In *Proceedings of the 19th International Society for Music Information Retrieval Conference (ISMIR 2018)*.

[4] Dong, H.-W., et al. 2021.
Towards Automatic Instrumentation by Learning to Separate Instruments in Music Score (ISMIR 2021 Paper #000019). In: *Proceedings of the 22nd International Society for Music Information Retrieval Conference (ISMIR 2021)*, edited by Lee et al., Online, Nov 7–12, 2021.

[5] Nabeoka, Takuto, Eita Nakamura & Kazuyoshi Yoshii. 2023.
Automatic Orchestration of Piano Scores for Wind Bands with User-Specified Instrumentation.
In *Proceedings of the 2023 International Conference on Computational Music Modeling and Retrieval (CMMR 2023)*.

[6] Le, Dinh-Viet-Toan & Yi-Hsuan Yang. 2025.
METEOR: Melody-aware Texture-controllable Symbolic Music Re-Orchestration via Transformer VAE. In: *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI-25)*, Montreal, Canada, 10126–10134. DOI:10.24963/ijcai.2025/1125.

[7] Crestel, L., & Esling, P 2017.
A Database Linking Piano and Orchestral MIDI for Music Analysis and Generation.